
MONTRÉAL.AI ACADEMY: ARTIFICIAL INTELLIGENCE 101

FIRST WORLD-CLASS OVERVIEW OF AI FOR ALL

VIP AI 101 CHEATSHEET

A PREPRINT

Vincent Boucher*
MONTRÉAL.AI
Montreal, Quebec, Canada
info@montreal.ai

July 6, 2019

ABSTRACT

For the purpose of entrusting all sentient beings with powerful AI tools to learn, deploy and scale AI in order to enhance their prosperity, to settle planetary-scale problems and to inspire those who, with AI, will shape the 21st Century, **MONTRÉAL.AI** introduces this **VIP AI 101 CheatSheet** for All.

Curated Open-Source Codes and Science: <http://www.academy.montreal.ai/>.

Keywords AI-First · Artificial Intelligence · Deep Learning · GANs · Intelligent Agent

1 AI-First

TODAY'S ARTIFICIAL INTELLIGENCE IS POWERFUL AND ACCESSIBLE TO ALL. AI opens up a world of new possibilities. To pioneer AI-First innovations advantages: start by exploring how to apply AI in ways never thought of.

"Breakthrough in machine learning would be worth 10 Microsofts." — Bill Gates

2 Getting Started

Tinker with neural networks in the browser with *TensorFlow Playground* <http://playground.tensorflow.org/>.

Papers With Code (*Learn Python 3 in Y minutes*²) <https://paperswithcode.com/state-of-the-art>.

2.1 In the Cloud

Colab³. Practice Immediately⁴. Labs⁵: Introduction to Deep Learning (MIT 6.S191)

- Free GPU compute via Colab <https://colab.research.google.com/notebooks/welcome.ipynb>.
- Six easy ways to run your Jupyter Notebook in the cloud⁶.

^{*}Founding Chairman at MONTRÉAL.AI <http://www.montreal.ai>.

²<https://learnxinyminutes.com/docs/python3/>

³<https://medium.com/tensorflow/colab-an-easy-way-to-learn-and-use-tensorflow-d74d1686e309>

⁴<https://colab.research.google.com/github/GokuMohandas/practicalAI/>

⁵https://colab.research.google.com/github/aamini/introtodeeplearning_labs

⁶<https://www.dataschool.io/cloud-services-for-jupyter-notebook/>

2.2 On a Local Machine

JupyterLab is an interactive development environment for working with notebooks, code and data⁷.

- Install Anaconda <https://www.anaconda.com/download/> and launch ‘Anaconda Navigator’
- Update Jupyterlab and launch the application. Under Notebook, click on ‘Python 3’

3 Deep Learning

Deep learning allows computational models that are composed of multiple processing layers to learn REPRESENTATIONS of (raw) data with multiple levels of abstraction[2]. At a high-level, neural networks are either encoders, decoders, or a combination of both⁸. Introductory course <http://introtodeeplearning.com>. See also Table 1.

“DL is essentially a new style of programming – “differentiable programming” – and the field is trying to work out the reusable constructs in this style. We have some: convolution, pooling, LSTM, GAN, VAE, memory units, routing units, etc.” — Thomas G. Dietterich

Table 1: Types of Learning, by Alex Graves at NeurIPS 2018

Name	With Teacher	Without Teacher
Active	<i>Reinforcement Learning / Active Learning</i>	<i>Intrinsic Motivation / Exploration</i>
Passive	<i>Supervised Learning</i>	<i>Unsupervised Learning</i>

“If you have a large big dataset and you train a very big neural network, then success is guaranteed!” — Ilya Sutskever

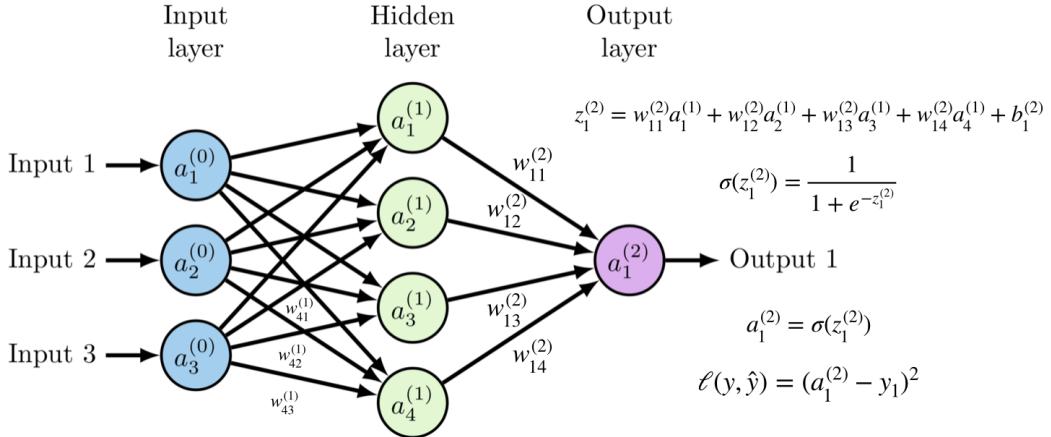


Figure 1: Multilayer perceptron (MLP).

“When you first study a field, it seems like you have to memorize a zillion things. You don’t. What you need is to identify the 3-5 core principles that govern the field. The million things you thought you had to memorize are various combinations of the core principles.” — J. Reed

- 1. Multiply things together
 - 2. Add them up
 - 3. Replaces negatives with zeros
 - 4. Return to step 1, a hundred times.”
- Jeremy Howard

⁷<https://blog.jupyter.org/jupyterlab-is-ready-for-users-5a6f039b8906>

⁸<https://github.com/lexfridman/mit-deep-learning>

Deep learning (*distributed representations + composition*) is a general-purpose learning procedure.

- ❖ Dive into Deep Learning <http://d2l.ai>.
- ❖ Minicourse in Deep Learning with PyTorch⁹.
- ❖ Deep Learning. The full deck of (600+) slides, by Gilles Louppe¹⁰.
- ❖ A Selective Overview of Deep Learning <https://arxiv.org/abs/1904.05526>.
- ❖ A Recipe for Training Neural Networks <https://karpathy.github.io/2019/04/25/recipe/>.
- ❖ How to Choose Your First AI Project <https://hbr.org/2019/02/how-to-choose-your-first-ai-project>.
- ❖ Blog | MIT 6.S191 <https://medium.com/tensorflow/mit-introduction-to-deep-learning-4a6f8dde1f0c>.

3.1 Universal Approximation Theorem

Neural Networks + Gradient Descent + GPU¹¹:

- Infinitely flexible function: *Neural Network* (multiple hidden layers: Deep Learning)¹².
- All-purpose parameter fitting: *Backpropagation*^{13 14}.

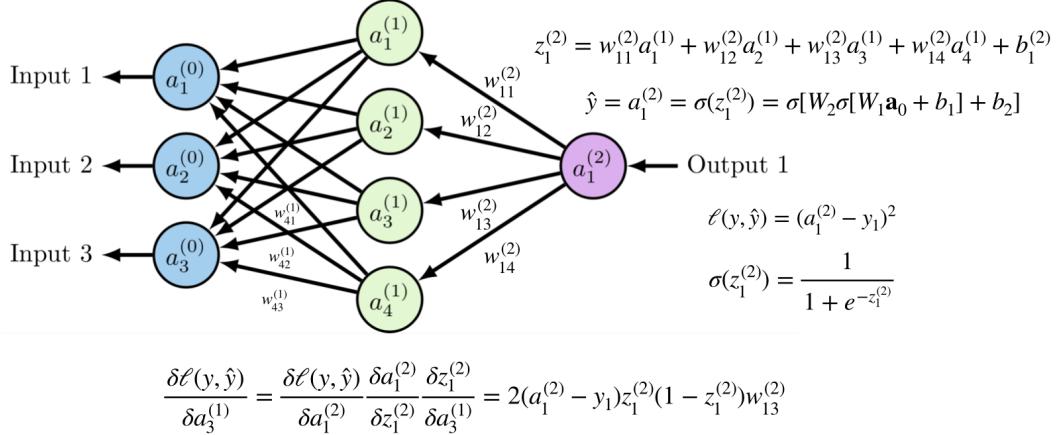


Figure 2: All-purpose parameter fitting: Backpropagation.

- Fast and scalable: *GPU*.

When a choice must be made, just feed the (raw) data to a deep neural network (Universal function approximators).

3.2 Convolution Neural Networks (Useful for Images | Space)

The deep convolutional network, inspired by Hubel and Wiesel's seminal work on early visual cortex, uses hierarchical layers of tiled convolutional filters to mimic the effects of receptive fields, thereby exploiting the local spatial correlations present in images[1]. See Figure 4. Demo <https://ml4a.github.io/demos/convolution/>.

A ConvNet is made up of Layers. Every Layer has a simple API: It transforms an input 3D volume to an output 3D volume with some differentiable function that may or may not have parameters¹⁵. Reading¹⁶.

In images, local combinations of edges form motifs, motifs assemble into parts, and parts form objects^{17 18}.

- ❖ TensorSpace (<https://tensorspace.org>) offers interactive 3D visualizations of *LeNet*, *AlexNet* and *Inceptionv3*.

⁹<https://github.com/Atcold/pytorch-Deep-Learning-Minicourse>

¹⁰<https://glouppe.github.io/info8010-deep-learning/pdf/lec-all.pdf>

¹¹http://wiki.fast.ai/index.php/Lesson_1_Notes

¹²<http://neuralnetworksanddeeplearning.com/chap4.html>

¹³https://github.com/DebPanigrahi/Machine-Learning/blob/master/back_prop.ipynb

¹⁴<https://www.jeremyjordan.me/neural-networks-training/>

¹⁵<http://cs231n.github.io/convolutional-networks/>

¹⁶<https://ml4a.github.io/ml4a/convnets/>

¹⁷<http://yosinski.com/deepvis>

¹⁸<https://distill.pub/2017/feature-visualization/>

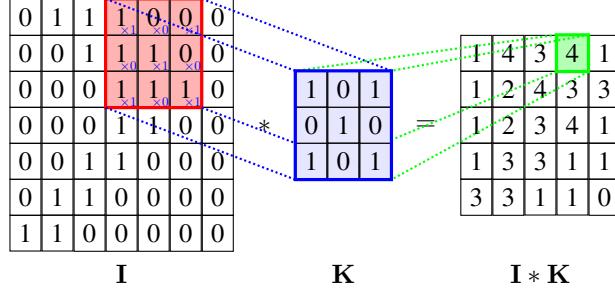
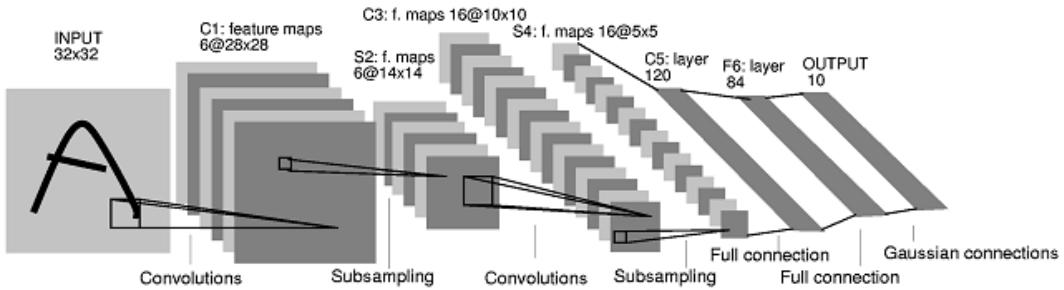
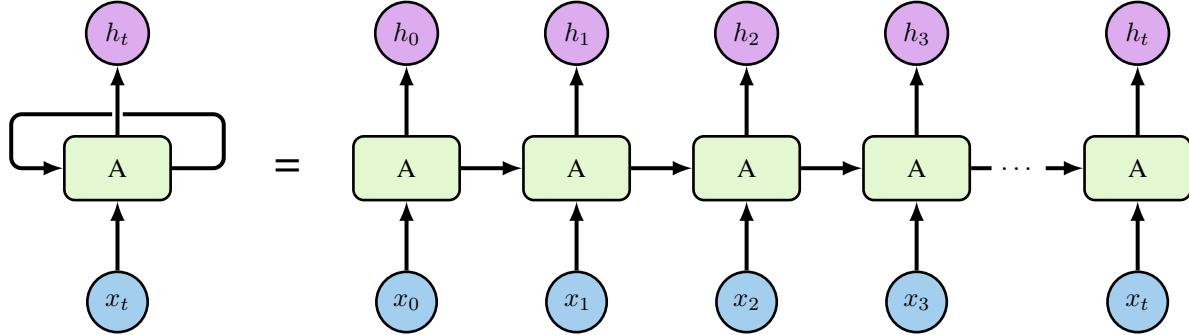
Figure 3: **2D Convolution**. Source: Cambridge Coding Academy

Figure 4: Architecture of LeNet-5, a Convolutional Neural Network. LeCun et al., 1998

3.3 Recurrent Neural Networks (Useful for Sequences | Time)

Recurrent neural networks are networks with loops in them, allowing information to persist¹⁹. RNNs process an input sequence one element at a time, maintaining in their hidden units a ‘state vector’ that implicitly contains information about the history of all the past elements of the sequence[2]. For sequential inputs. See Figure 6.

Figure 5: **RNN Layers Reuse Weights for Multiple Timesteps**.

"I feel like a significant percentage of Deep Learning breakthroughs ask the question "how can I reuse weights in multiple places?" – Recurrent (LSTM) layers reuse for multiple timesteps – Convolutional layers reuse in multiple locations. – Capsules reuse across orientation." — Andrew Trask

- ❖ Long Short-Term-Memory (LSTM), Sepp Hochreiter and Jürgen Schmidhuber²⁰.
- ❖ Can Neural Networks Remember? Slides by Vishal Gupta: http://vishalgupta.me/deck/char_lstm/.
- ❖ Understanding LSTM Networks <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>.
- ❖ The Unreasonable Effectiveness of Recurrent Neural Networks, blog (2015) by Andrej Karpathy²¹.

¹⁹<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

²⁰<https://www.bioinf.jku.at/publications/older/2604.pdf>

²¹<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

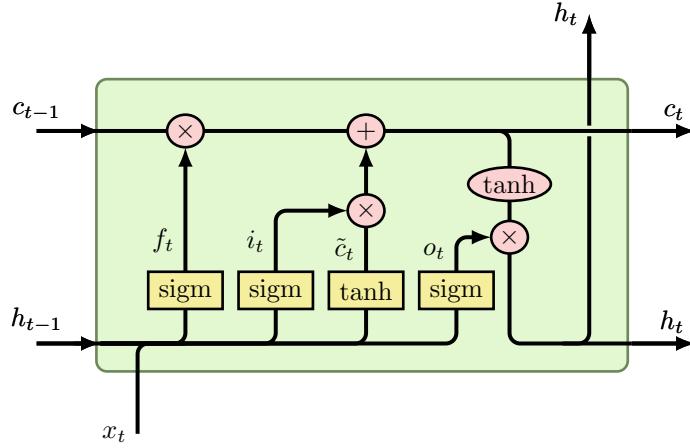


Figure 6: "Long Short-Term-Memory" (LSTM) Cell.

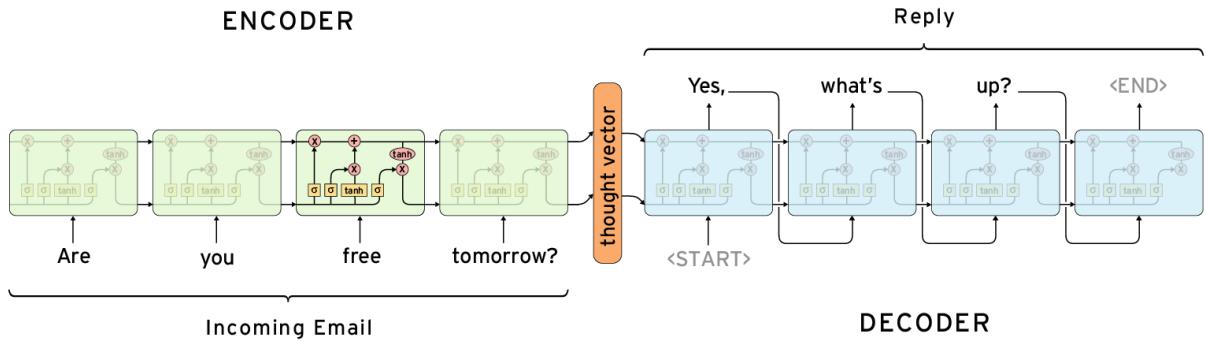


Figure 7: Google Smart Reply System is built on a pair of recurrent neural networks. Diagram by Chris Olah

- ❖ Attention and Augmented Recurrent Neural Networks <https://distill.pub/2016/augmented-rnns/>.
- ❖ Attention Is All You Need, Vaswani et al. <https://arxiv.org/abs/1706.03762>.
- ❖ Transformer model for language understanding. Tutorial showing how to write Transformer in TensorFlow 2.0²².

3.4 Unsupervised Learning

True intelligence will require independent learning strategies.

Unsupervised learning is a paradigm for creating AI that learns without a particular task in mind: learning for the sake of learning²³. It captures some characteristics of the joint distribution of the observed random variables (learn the underlying structure). The variety of tasks include density estimation, dimensionality reduction, and clustering.[4]²⁴.

Self-supervised learning is derived from unsupervised learning where the data provides the supervision. E.g. Word2vec²⁵, a technique for learning vector representations of words, or word **embeddings**. An embedding is a mapping from discrete objects, such as words, to vectors of real numbers²⁶.

²²<https://www.tensorflow.org/alpha/tutorials/sequences/transformer>

²³<https://deepmind.com/blog/unsupervised-learning/>

²⁴https://media.neurips.cc/Conferences/NIPS2018/Slides/Deep_Unsupervised_Learning.pdf

²⁵<https://jalammar.github.io/illustrated-word2vec/>

²⁶<http://projector.tensorflow.org>

3.4.1 Generative Adversarial Networks

Simultaneously train two models: a generative model G that captures the data distribution, and a discriminative model D that estimates the probability that a sample came from the training data rather than G. The training procedure for G is to maximize the probability of D making a mistake. This framework corresponds to a minimax two-player game[3].

$$\min_{\theta_g} \max_{\theta_d} [\mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D_{\theta_d}(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D_{\theta_d}(G_{\theta_g}(z)))]] \quad (1)$$

"What I cannot create, I do not understand." — Richard Feynman

Goodfellow et al. used an interesting analogy where the generative model can be thought of as analogous to a team of counterfeiters, trying to produce fake currency and use it without detection, while the discriminative model is analogous to the police, trying to detect the counterfeit currency. Competition in this game drives both teams to improve their methods until the counterfeits are indistinguishable from the genuine articles. See Figure 8.

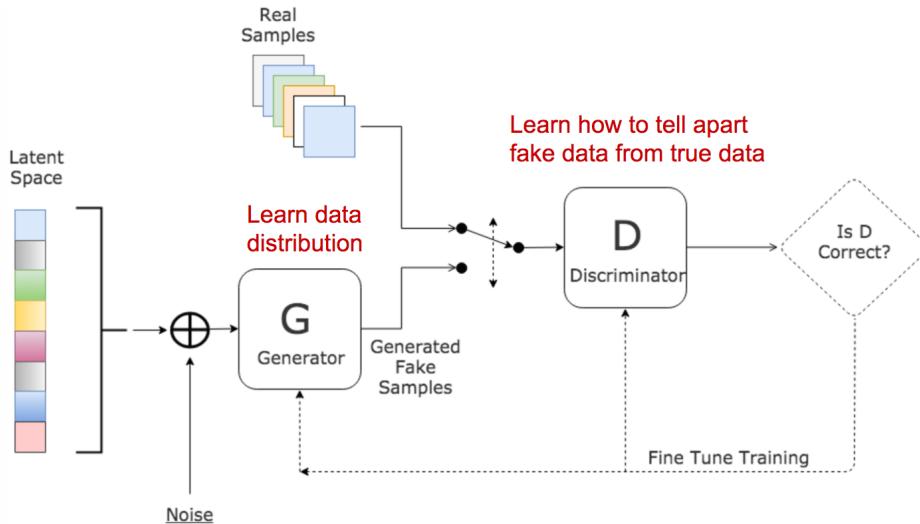


Figure 8: GAN: Neural Networks Architecture Pioneered by Ian Goodfellow at University of Montreal (2014).

StyleGAN: A Style-Based Generator Architecture for Generative Adversarial Networks

- Paper <http://stylegan.xyz/paper> | Code <https://github.com/NVlabs/stylegan>.
 - **StyleGAN for art.** Colab <https://colab.research.google.com/github/ak9250/stylegan-art>.
 - This Person Does Not Exist <https://thispersondoesnotexist.com>.
 - Which Person Is Real? <http://www.whichfaceisreal.com>.
 - This Resume Does Not Exist <https://thisresumedoestnotexist.com>.
 - This Waifu Does Not Exist <https://www.thiswaifudoesnotexist.net>.
 - Encoder for Official TensorFlow Implementation <https://github.com/Puzer/stylegan-encoder>.
 - How to recognize fake AI-generated images. By Kyle McDonald²⁷.
- ❖ Few-Shot Adversarial Learning of Realistic Neural Talking Head Models²⁸.
- ❖ Wasserstein GAN <http://www.depthfirstlearning.com/2019/WassersteinGAN>.
- ❖ GANSynth: Generate high-fidelity audio with GANs! Colab <http://goo.gl/magenta/gansynth-demo>.
- ❖ SC-FEGAN: Face Editing Generative Adversarial Network <https://github.com/JoYoungjoo/SC-FEGAN>.
- ❖ CariGANs: Unpaired Photo-to-Caricature Translation. Cao et al.: <https://cari-gan.github.io>.
- ❖ GANpaint Paint with GAN units <http://gandissect.res.ibm.com/ganpaint.html>.
- ❖ PyTorch pretrained BigGAN <https://github.com/huggingface/pytorch-pretrained-BigGAN>.

²⁷<https://medium.com/@kcimc/how-to-recognize-fake-ai-generated-images-4d1f6f9a2842>

²⁸<https://arxiv.org/abs/1905.08233>

- ❖ Demo of BigGAN in an official Colaboratory notebook (backed by a GPU) https://colab.research.google.com/github/tensorflow/hub/blob/master/examples/colab/biggan_generation_with_tf_hub.ipynb

3.4.2 Variational AutoEncoder

Variational Auto-Encoders (VAEs) are powerful models for learning low-dimensional representations. See Figure 9. Disentangled representations are defined as ones where a change in a single unit of the representation corresponds to a change in single factor of variation of the data while being invariant to others (Bengio et al. (2013)).

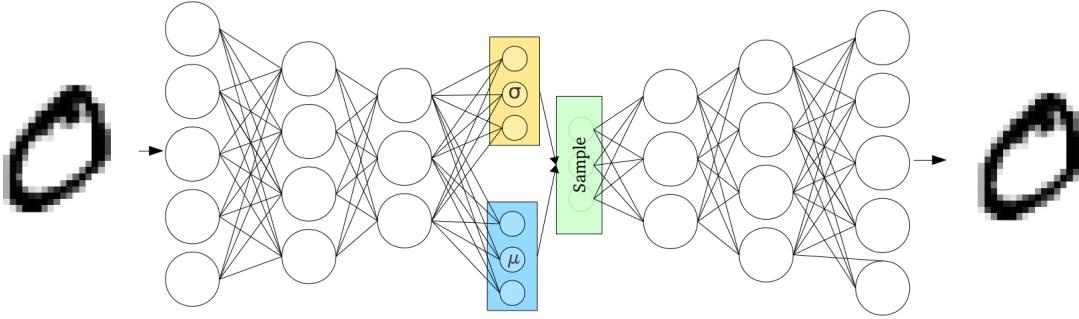


Figure 9: Variational Autoencoders (VAEs): Powerful Generative Models.

- ❖ Colab²⁹: "Debiasing Facial Detection Systems." [AIEthics](#)
- ❖ SpaceSheet: Interactive Latent Space Exploration with a Spreadsheet <https://vusd.github.io/spacesheet/>.
- ❖ MusicVAE: Learning latent spaces for musical scores <https://magenta.tensorflow.org/music-vae>.
- ❖ Slides: A Few Unusual Autoencoders <https://colinraffel.com/talks/vector2018few.pdf>.
- ❖ Generative models in Tensorflow 2 <https://github.com/timsainb/tensorflow2-generative-models/>.
- ❖ Reading: Disentangled VAE's (DeepMind 2016) <https://arxiv.org/abs/1606.05579>.

3.4.3 Natural Language Processing (NLP) | BERT: A New Era in NLP

BERT (Bidirectional Encoder Representations from Transformers)[6] is a *deeply bidirectional, unsupervised language representation*, pre-trained using only a plain text corpus (in this case, Wikipedia)³⁰.

- Reading: Unsupervised pre-training of an LSTM followed by supervised fine-tuning[7].
- TensorFlow code and pre-trained models for BERT <https://github.com/google-research/bert>.
- Better Language Models and Their Implications³¹.

"I think transfer learning is the key to general intelligence. And I think the key to doing transfer learning will be the acquisition of conceptual knowledge that is abstracted away from perceptual details of where you learned it from." — Demis Hassabis

- ❖ How to Build OpenAI's GPT-2: "*The AI That's Too Dangerous to Release*"³².
- ❖ Play with BERT with your own data using TensorFlow Hub https://colab.research.google.com/github/google-research/bert/blob/master/predicting_movie_reviews_with_bert_on_tf_hub.ipynb.

²⁹https://colab.research.google.com/github/aamini/introtodeeplearning_labs/blob/master/lab2/Part2_debiasing_solution.ipynb

³⁰<https://ai.googleblog.com/2018/11/open-sourcing-bert-state-of-art-pre.html>

³¹<https://blog.openai.com/better-language-models/>

³²<https://blog.floydhub.com/gpt2/>

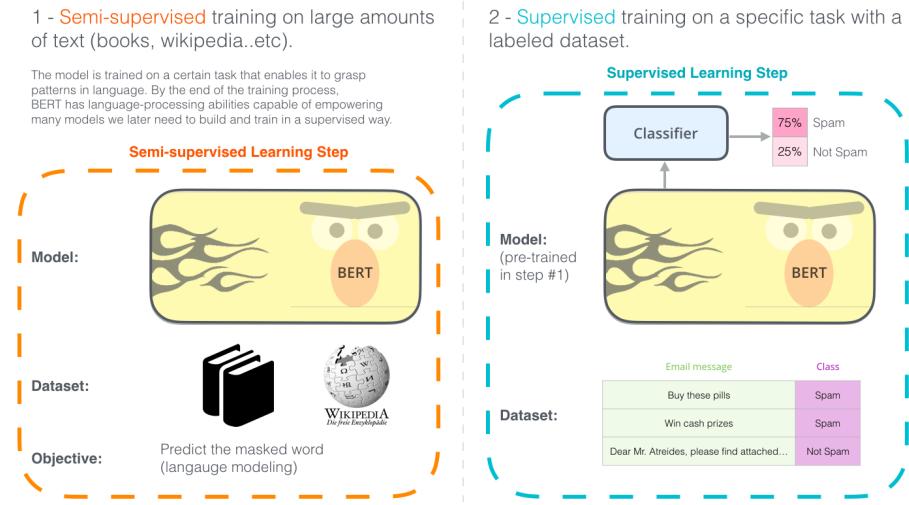


Figure 10: The two steps of how BERT is developed. Source <https://jalammar.github.io/illustrated-bert/>.

4 Autonomous Agents

An **autonomous agent** is any device that perceives its environment and takes actions that maximize its chance of success at some goal. At the bleeding edge of AI, autonomous agents can learn from experience, simulate worlds and orchestrate meta-solutions. Here's an informal definition³³ of the *universal intelligence* of agent π ³⁴:

$$\Upsilon(\pi) := \sum_{\mu \in E} 2^{-K(\mu)} V_\mu^\pi \quad (2)$$

"Intelligence measures an agent's ability to achieve goals in a wide range of environments." — Shane Legg

4.1 Deep Reinforcement Learning

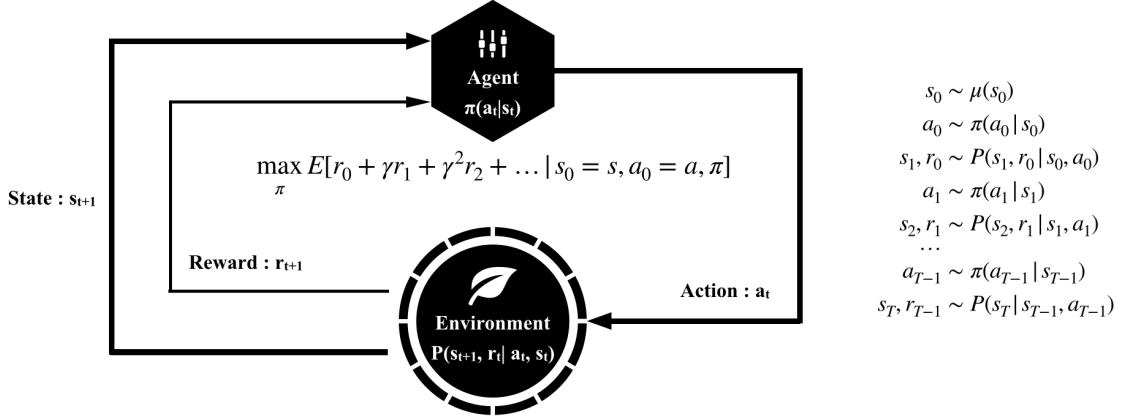


Figure 11: An Agent Interacts with an Environment.

Reinforcement learning (RL) studies how an agent can learn how to achieve goals in a complex, uncertain environment (Figure 12) [5]. Recent superhuman results in many difficult environments combine deep learning with RL (*Deep Reinforcement Learning*). See Figure 12 for a taxonomy of RL algorithms.

³³<https://arxiv.org/abs/0712.3329>

³⁴Where μ is an environment, K is the Kolmogorov complexity function, E is the space of all computable reward summable environmental measures with respect to the reference machine U and the value function V_μ^π is the agent's "ability to achieve".

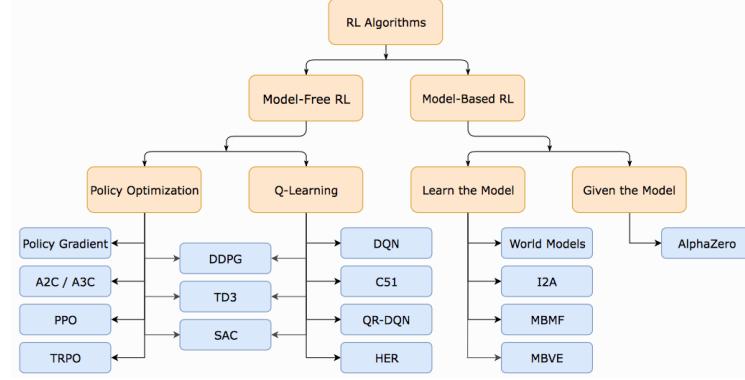


Figure 12: A Taxonomy of RL Algorithms. Source: Spinning Up in Deep RL by Achiam et al. | OpenAI

4.1.1 Model-Free RL | Value-Based

The goal in RL is to train the agent to maximize the discounted sum of all future rewards R_t , called the return:

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \quad (3)$$

The Q-function captures the expected total future reward an agent in state s can receive by executing a certain action a :

$$Q(s, a) = E[R_t] \quad (4)$$

The optimal policy should choose the action a that maximizes $Q(s, a)$:

$$\pi^*(s) = \text{argmax}_a Q(s, a) \quad (5)$$

- **Q-Learning:** Playing Atari with Deep Reinforcement Learning (DQN). Mnih et al, 2013[10].

TF-Agents (DQN Tutorial) | Colab <https://colab.research.google.com/github/tensorflow/agents>.

4.1.2 Model-Free RL | Policy-Based

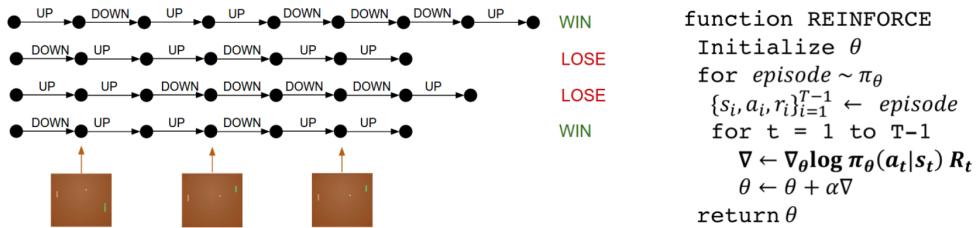


Figure 13: Policy Gradient Directly Optimizes the Policy.

Run a policy for a while (code: <https://gist.github.com/karpathy/a4166c7fe253700972fcbe77e4ea32c5>):

$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T) \quad (6)$$

Increase probability of actions that lead to high rewards and decrease probability of actions that lead to low rewards:

$$\nabla_\theta E_\tau[R(\tau)] = E_\tau \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi(a_t | s_t, \theta) R(\tau) \right] \quad (7)$$

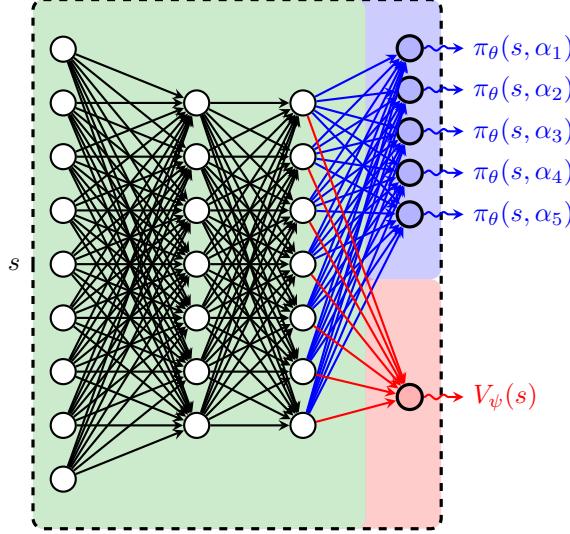


Figure 14: **Asynchronous Advantage Actor-Critic (A3C)**. Source: Petar Velickovic

- **Policy Optimization:** *Asynchronous Methods for Deep Reinforcement Learning* (A3C). Mnih et al, 2016[8].
- **Policy Optimization:** *Proximal Policy Optimization Algorithms* (PPO). Schulman et al, 2017[9].

4.1.3 Model-Based RL

In Model-Based RL, the agent generates predictions about the next state and reward before choosing each action.

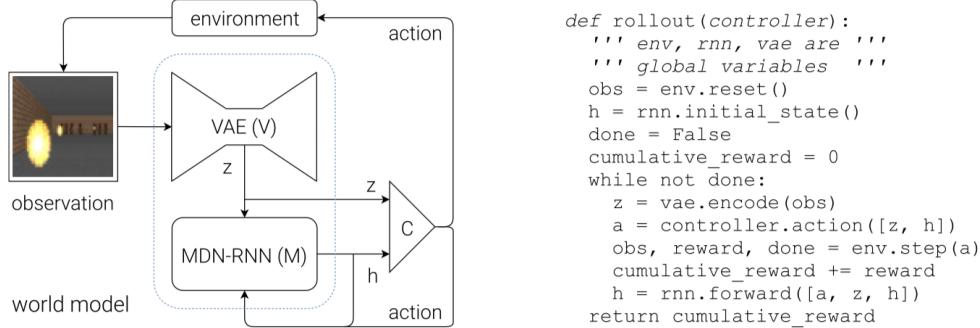


Figure 15: World Model's Agent consists of: Vision (V), Memory (M), and Controller (C). | Ha et al, 2018[11]

- **Learn the Model:** *Recurrent World Models Facilitate Policy Evolution* (World Models³⁵). The world model agent can be trained in an unsupervised manner to learn a compressed spatial and temporal representation of the environment. Then, a compact policy can be trained. See Figure 15. Ha et al, 2018[11].
- **Learn the Model:** *Learning Latent Dynamics for Planning from Pixels* <https://planetrl.github.io/>.
- **Given the Model:** *Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm* (AlphaZero). Silver et al, 2017[14]. AlphaGo Zero Explained In One Diagram³⁶.

³⁵<https://worldmodels.github.io>

³⁶https://applied-data.science/static/main/res/alpha_go_zero_cheat_sheet.png

4.1.4 Improving Agent Design

Via Reinforcement Learning: Blog³⁷. arXiv³⁸. ASTool <https://github.com/hardmaru/astool/>.
 Via Evolution: Video³⁹. Evolved Creatures <http://www.karlsims.com/evolved-virtual-creatures.html>.

4.1.5 OpenAI Baselines

High-quality implementations of reinforcement learning algorithms <https://github.com/openai/baselines>.
 Colab <https://colab.research.google.com/drive/1KKq9A3dRTq1q6bJmPyF0gg917gQyTjJI>.

4.1.6 Google Dopamine and A Zoo of Agents

Dopamine is a research framework for fast prototyping of reinforcement learning algorithms.⁴⁰

A Zoo of Atari-Playing Agents: Code⁴¹, Blog⁴² and Colaboratory notebook <https://colab.research.google.com/github/uber-research/atari-model-zoo/blob/master/colab/AtariZooColabDemo.ipynb>.

4.2 Evolution Strategies (ES)

Evolution and neural networks proved a potent combination in nature. Neuroevolution, which harnesses evolutionary algorithms to optimize neural networks, enables capabilities that are typically unavailable to gradient-based approaches, including learning neural network building blocks, architectures and even the algorithms for learning[12].

"... evolution — whether biological or computational — is inherently creative, and should routinely be expected to surprise, delight, and even outwit us." — The Surprising Creativity of Digital Evolution, Lehman et al.[21]

Neural architecture search has advanced to the point where it can outperform human-designed models[13].

Natural evolutionary strategy directly evolves the weights of a DNN and performs competitively with the best deep reinforcement learning algorithms, including deep Q-networks (DQN) and policy gradient methods (A3C)[20].

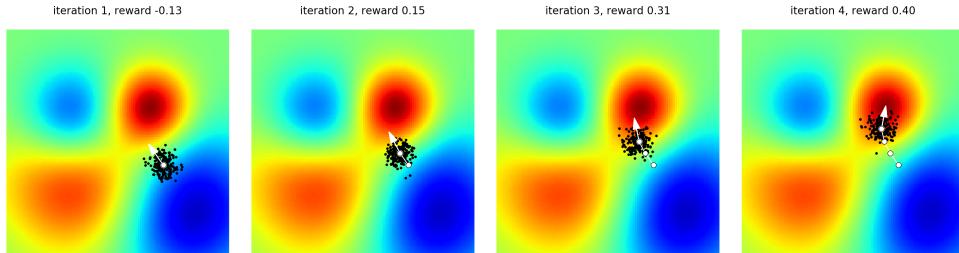


Figure 16: <https://colab.research.google.com/github/karpathy/randomfun/blob/master/es.ipynb>.

The ES algorithm is a “guess and check” process, where we start with some random parameters and then repeatedly:

1. Tweak the guess a bit randomly, and
2. Move our guess slightly towards whatever tweaks worked better.

"Evolution is a slow learning algorithm that with the sufficient amount of compute produces a human brain." — Wojciech Zaremba

³⁷<https://designrl.github.io>

³⁸<https://arxiv.org/abs/1810.03779>

³⁹https://youtu.be/JBgG_VSP7f8

⁴⁰<https://github.com/google/dopamine>

⁴¹<https://github.com/uber-research/atari-model-zoo>

⁴²<https://eng.uber.com/atari-zoo-deep-reinforcement-learning/>

Demos: ES on CartPole-v1⁴³ and ES on LunarLanderContinuous-v2⁴⁴.

VAE+CPPN+GAN https://colab.research.google.com/drive/1_OoZ3z_C5Jl5gnxD0E9VEMCTs-F18pvM.
A Visual Guide to ES <http://blog.otoro.net/2017/10/29/visual-evolution-strategies/>.

4.3 Self Play

Silver et al.[15] introduced an algorithm based solely on reinforcement learning, without human data, guidance or domain knowledge. Starting tabula rasa (and being its own teacher!), AlphaGo Zero achieved superhuman performance. AlphaGo Zero showed that algorithms matter much more than big data and massive amounts of computation.

"Self-Play is Automated Knowledge Creation." — Carlos E. Perez

Self-play mirrors similar insights from coevolution. Transfer learning is the key to go from self-play to the real world⁴⁵.

ALPHAGO ZERO CHEAT SHEET

The training pipeline for AlphaGo Zero consists of three stages, executed in parallel

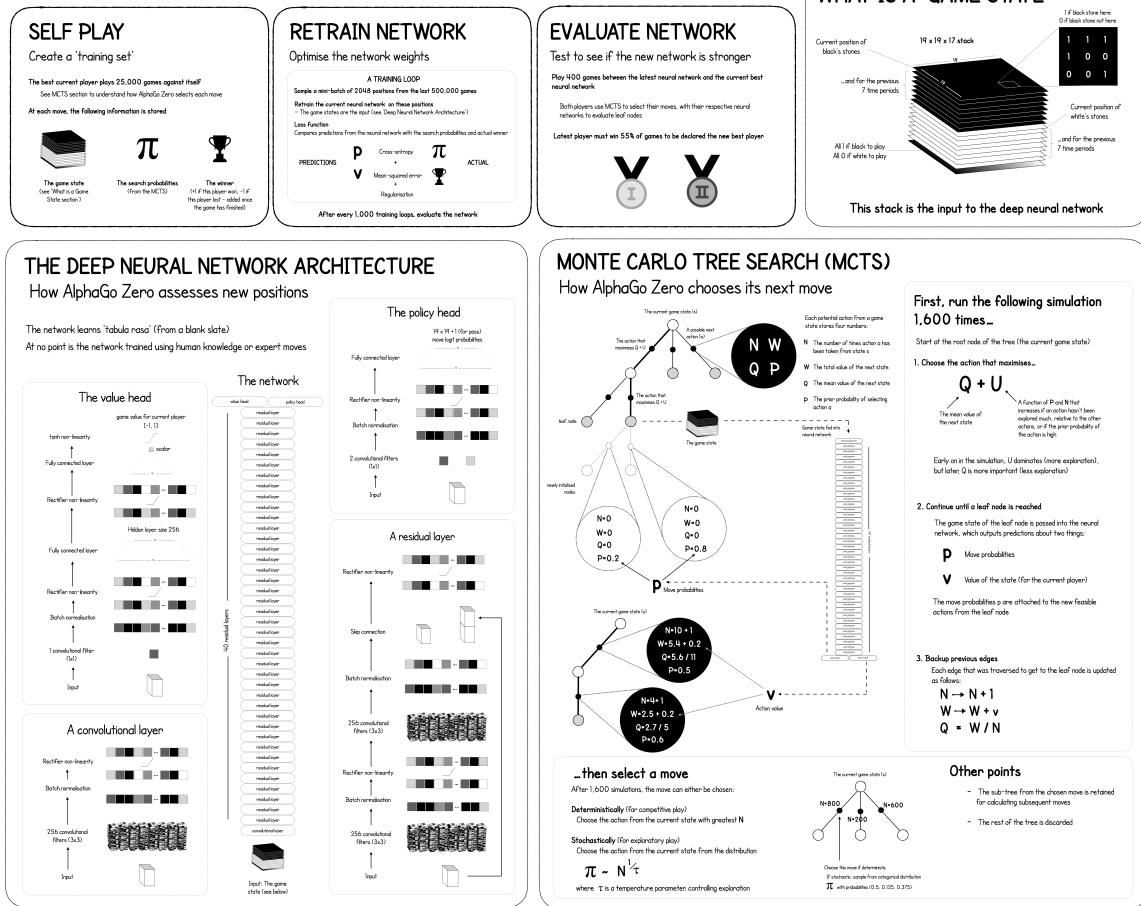


Figure 17: Ref.: https://applied-data.science/static/main/res/alpha_go_zero_cheat_sheet.png.

"Open-ended self play produces: Theory of mind, negotiation, social skills, empathy, real language understanding." — Ilya Sutskever, Meta Learning and Self Play

⁴³<https://colab.research.google.com/drive/1bMZWHDhm-mT9NJENWoVewUks7cGV10go>

⁴⁴https://colab.research.google.com/drive/1lvYKjFtc_C_8njCKD-MnXEW8LPS2RPr6

⁴⁵<http://metalearning-symposium.ml>

TensorFlow.js Implementation of DeepMind's AlphaZero Algorithm for Chess. Live Demo⁴⁶ | Code⁴⁷

An open-source implementation of the AlphaGoZero algorithm <https://github.com/tensorflow/minigo>
 ELF OpenGo: An Open Reimplementation of AlphaZero, Tian et al.: <https://arxiv.org/abs/1902.04522>.

4.4 Deep Meta-Learning

Learning to Learn[16]. A meta-learning algorithm takes in a distribution of tasks, where each task is a learning problem, and it produces a quick learner — a learner that can generalize from a small number of examples[17].

"The notion of a neural "architecture" is going to disappear thanks to meta learning." — Andrew Trask

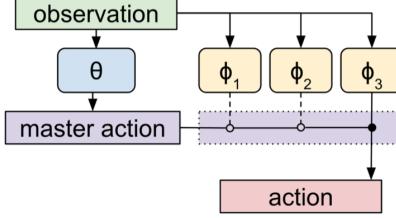


Figure 18: Meta Learning Shared Hierarchies[18] (The Lead Author is in High School!)

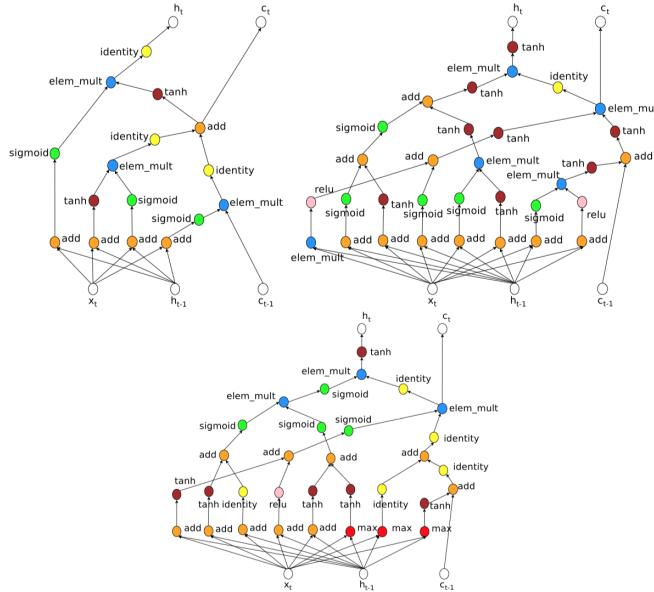


Figure 19: A comparison of the original LSTM cell vs. two new good generated. Top left: LSTM cell. [19]

"The future of high-level APIs for AI is... a problem-specification API. Currently we only search over network weights, thus "problem specification" involves specifying a model architecture. In the future, it will just be: "tell me what data you have and what you are optimizing"." — François Chollet

"Causal Reasoning from Meta-reinforcement Learning," Dasgupta et al.: <https://arxiv.org/abs/1901.08162>

Colaboratory reimplementation of MAML (Model-Agnostic Meta-Learning) in TF 2.0⁴⁸

⁴⁶<https://frpayes.github.io/lc0-js/engine.html>

⁴⁷<https://github.com/frpayes/lc0-js/>

⁴⁸<https://colab.research.google.com/github/mari-linhaires/tensorflow-maml/blob/master/maml.ipynb>

4.5 Multi-Agent Populations

"We design a Theory of Mind neural network – a ToMnet – which uses meta-learning to build models of the agents it encounters, from observations of their behaviour alone." — Machine Theory of Mind, Rabinowitz et al.[24]

Cooperative Agents. Learning to Model Other Minds, by OpenAI[23], is an algorithm which accounts for the fact that other agents are learning too, and discovers self-interested yet collaborative strategies. Also: OpenAI Five⁴⁹.

"Artificial Intelligence is about recognising patterns, Artificial Life is about creating patterns." — Mizuki Oka et al.

Active Learning Without Teacher. In *Intrinsic Social Motivation via Causal Influence in Multi-Agent RL*, Jaques et al. (2018) <https://arxiv.org/abs/1810.08647> propose an intrinsic reward function designed for multi-agent RL (MARL), which awards agents for having a causal influence on other agents' actions. Open-source implementation⁵⁰.

"Open-ended Learning in Symmetric Zero-sum Games," Balduzzi et al.: <https://arxiv.org/abs/1901.08106>
Neural MMO: a massively multiagent env. for simulations with many long-lived agents. Code⁵¹ and 3D Client⁵².

5 Environments

Platforms for training autonomous agents.

"Situation awareness is the perception of the elements in the environment within a volume of time and space, and the comprehension of their meaning, and the projection of their status in the near future." — Endsley (1987)

5.1 OpenAI Gym

The OpenAI Gym <https://gym.openai.com/> (Blog⁵³ | GitHub⁵⁴) is a toolkit for developing and comparing reinforcement learning algorithms. What makes the gym so great is a common API around environments.

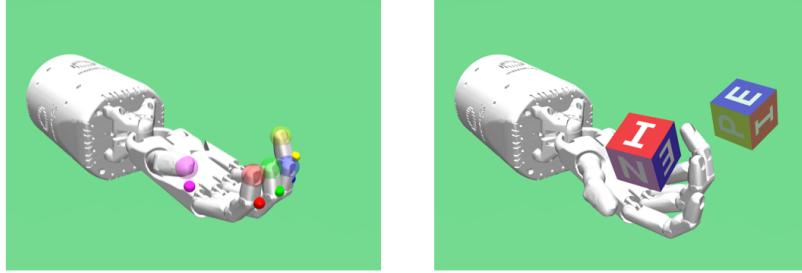


Figure 20: Robotics Environments <https://blog.openai.com/ingredients-for-robotics-research/>

How to create new environments for Gym⁵⁵. Examples: OpenAI Gym Environment for Trading⁵⁶.

5.2 DeepMind Lab

DeepMind Lab: A customisable 3D platform for agent-based AI research <https://github.com/deepmind/lab>.

- DeepMind Control Suite https://github.com/deepmind/dm_control.
- Convert DeepMind Control Suite to OpenAI Gym Envs <https://github.com/zuxingdong/dm2gym>.

⁴⁹<https://blog.openai.com/openai-five/>

⁵⁰https://github.com/eugenevinitksy/sequential_social_dilemma_games

⁵¹<https://github.com/openai/neural-mmo>

⁵²<https://github.com/jsuarez5341/neural-mmo-client>

⁵³<https://blog.openai.com/openai-gym-beta/>

⁵⁴<https://github.com/openai/gym>

⁵⁵<https://github.com/openai/gym/blob/master/docs/creating-environments.md>

⁵⁶<https://github.com/hackthemarket/gym-trading>

5.3 Unity ML-Agents

Unity ML Agents allows to create environments where intelligent agents (*Single Agent, Cooperative and Competitive Multi-Agent* and *Ecosystem*) can be trained using RL, neuroevolution, or other ML methods <https://unity3d.ai>.

- Getting Started with Marathon Environments for Unity ML-Agents⁵⁷.
- Arena: A General Evaluation Platform and Building Toolkit for Multi-Agent Intelligence⁵⁸.

5.4 POET: Paired Open-Ended Trailblazer

Diversity is the premier product of evolution. Endlessly generate increasingly complex and diverse learning environments⁵⁹. Open-endedness could generate learning algorithms reaching human-level intelligence[22].

- Implementation of the POET algorithm <https://github.com/uber-research/poet>.
- ❖ AI-GAs: AI-generating algorithms, an alternate paradigm for producing general artificial intelligence⁶⁰.

6 Datasets

Google Dataset Search Beta (Blog⁶¹) <https://toolbox.google.com/datasetsearch>.

TensorFlow Datasets: load a variety of public datasets into TensorFlow programs (Blog⁶² | Colab⁶³).

7 Deep-Learning Hardware

A Full Hardware Guide to Deep Learning, by Tim Dettmers⁶⁴.

Which GPU(s) to Get for Deep Learning, by Tim Dettmers⁶⁵.



Figure 21: Edge TPU - Dev Board <https://coral.withgoogle.com/products/dev-board/>

Build AI that works offline with Coral Dev Board, Edge TPU, and TensorFlow Lite, by Daniel Situnayake⁶⁶. Jetson Nano. A small but mighty AI computer to create intelligent systems⁶⁷.

⁵⁷<https://towardsdatascience.com/gettingstartedwithmarathonenvs-v0-5-0a-c1054a0b540c>

⁵⁸<https://arxiv.org/abs/1905.08085>

⁵⁹<https://eng.uber.com/poet-open-ended-deep-learning/>

⁶⁰<https://arxiv.org/abs/1905.10985>

⁶¹<https://www.blog.google/products/search/making-it-easier-discover-datasets/>

⁶²<https://medium.com/tensorflow/introducing-tensorflow-datasets-c7f01f7e19f3>

⁶³<https://colab.research.google.com/github/tensorflow/datasets/blob/master/docs/overview.ipynb>

⁶⁴<http://timdettmers.com/2018/12/16/deep-learning-hardware-guide/>

⁶⁵<http://timdettmers.com/2019/04/03/which-gpu-for-deep-learning/>

⁶⁶<https://medium.com/tensorflow/build-ai-that-works-offline-with-coral-dev-board-edge-tpu-and-tensorflow-lite-70>

⁶⁷<https://www.nvidia.com/en-us/autonomous-machines/embedded-systems/jetson-nano/>

8 Deep-Learning Software

TensorFlow

- tf.keras (TensorFlow 2.0) for Researchers: Crash Course. Colab⁶⁸.
- TensorFlow 2.0: basic ops, gradients, data preprocessing and augmentation, training and saving. Colab⁶⁹.
- TensorBoard in Jupyter Notebooks. Colab⁷⁰.
- TensorFlow Lite for Microcontrollers⁷¹.

PyTorch

- PyTorch primer. Colab⁷².
- PyTorch internals <http://blog.ezyang.com/2019/05/pytorch-internals/>

9 AI Art | A New Day Has Come in Art Industry



Figure 22: On October 25, 2018, the first AI artwork ever sold at Christie’s auction house fetched USD 432,500.

The code (*art-DCGAN*) for the first artificial intelligence artwork ever sold at Christie’s auction house (Figure 22) is a modified implementation of DCGAN focused on generative art: <https://github.com/robbiebarrat/art-dcgan>.

- **TensorFlow Magenta**. An open source research project exploring the role of ML in the creative process.⁷³.
- **Magenta Studio**. A suite of free music-making tools using machine learning models!⁷⁴.
- **Style Transfer Tutorial** https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/generative/style_transfer.ipynb
- **AI x AR Paper Cubes** <https://experiments.withgoogle.com/paper-cubes>.
- **Photo Wake-Up** <https://grail.cs.washington.edu/projects/wakeup/>.
- **COLLECTION**. AI Experiments <https://experiments.withgoogle.com/ai>.

"The Artists Creating with AI Won't Follow Trends; THEY WILL SET THEM." — The House of Montréal.AI Fine Arts

MuseNet. Generate Music Using Many Different Instruments and Styles!⁷⁵.
 Tuning Recurrent Neural Networks with Reinforcement Learning⁷⁶.

⁶⁸<https://colab.research.google.com/drive/14CvUNTaX10FHDfaKaaZzrBsvMfhCOHIR>

⁶⁹https://colab.research.google.com/github/zaidalyafeai/Notebooks/blob/master/TF_2_0.ipynb

⁷⁰https://colab.research.google.com/github/tensorflow/tensorboard/blob/master/docs/r2/get_started.ipynb

⁷¹<https://petewarden.com/2019/03/07/launching-tensorflow-lite-for-microcontrollers/>

⁷²<https://colab.research.google.com/drive/1DgkVmi6GksW0ByhYVQpyUB4Rk3PUq0Cp>

⁷³<https://magenta.tensorflow.org>

⁷⁴<https://magenta.tensorflow.org/studio>

⁷⁵<https://openai.com/blog/musenet/>

⁷⁶<https://magenta.tensorflow.org/2016/11/09/tuning-recurrent-networks-with-reinforcement-learning>

Discovering Visual Patterns in Art Collections with Spatially-consistent Feature Learning. Shen et al.⁷⁷. Deep Multispectral Painting Reproduction via Multi-Layer, Custom-Ink Printing. Shi et al.⁷⁸.

10 AI Macrostrategy: Aligning AGI with Human Interests

Montréal.AI Governance: Policies at the intersection of AI, Ethics and Governance.

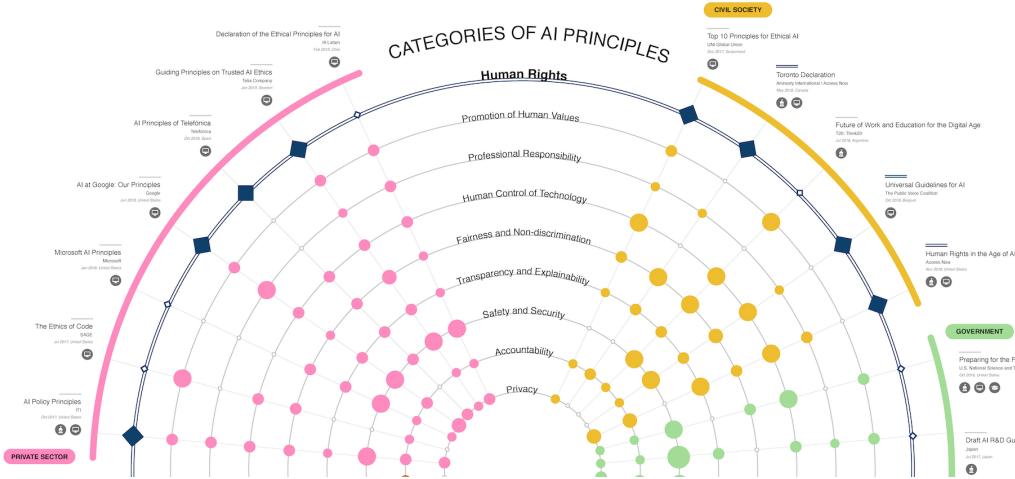


Figure 23: A Map of Ethical and Right-Based Approaches <https://ai-hr.cyber.harvard.edu/primp-viz.html>

- ❖ **AI Index.** <http://aiindex.org>.
- ❖ **Malicious AI Report.** <https://arxiv.org/pdf/1802.07228.pdf>.
- ❖ **Artificial Intelligence and Human Rights.** <https://ai-hr.cyber.harvard.edu>.

"(AI) will rank among our greatest technological achievements, and everyone deserves to play a role in shaping it." — Fei-Fei Li

References

- [1] Mnih et al. Human-Level Control Through Deep Reinforcement Learning. In *Nature* 518, pages 529–533. 26 February 2015. <https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf>
- [2] Yann LeCun, Yoshua Bengio and Geoffrey Hinton. Deep Learning. In *Nature* 521, pages 436–444. 28 May 2015. <https://www.cs.toronto.edu/~hinton/absps/NatureDeepReview.pdf>
- [3] Goodfellow et al. Generative Adversarial Networks. *arXiv preprint arXiv:1406.2661*, 2014. <https://arxiv.org/abs/1406.2661>
- [4] Yoshua Bengio, Andrea Lodi, Antoine Prouvost. Machine Learning for Combinatorial Optimization: a Methodological Tour d’Horizon. *arXiv preprint arXiv:1811.06128*, 2018. <https://arxiv.org/abs/1811.06128>
- [5] Brockman et al. OpenAI Gym. 2016. <https://gym.openai.com>
- [6] Devlin et al. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *arXiv preprint arXiv:1810.04805*, 2018. <https://arxiv.org/abs/1810.04805>
- [7] Dai et al. Semi-supervised Sequence Learning. *arXiv preprint arXiv:1511.01432*, 2015. <https://arxiv.org/abs/1511.01432>
- [8] Mnih et al. Asynchronous Methods for Deep Reinforcement Learning. *arXiv preprint arXiv:1602.01783*, 2016. <https://arxiv.org/abs/1602.01783>
- [9] Schulman et al. Proximal Policy Optimization Algorithms. *arXiv preprint arXiv:1707.06347*, 2017. <https://arxiv.org/abs/1707.06347>

⁷⁷<https://arxiv.org/pdf/1903.02678.pdf>

⁷⁸<http://people.csail.mit.edu/liangs/papers/ToG18.pdf>

- [10] Mnih et al. Playing Atari with Deep Reinforcement Learning. *DeepMind Technologies*, 2013. <https://www.cs.toronto.edu/~vmnih/docs/dqn.pdf>
- [11] Ha et al. Recurrent World Models Facilitate Policy Evolution. *arXiv preprint arXiv:1809.01999*, 2018. <https://arxiv.org/abs/1809.01999>
- [12] Kenneth et al. Designing neural networks through neuroevolution. In *Nature Machine Intelligence* VOL 1, pages 24–35. January 2019. <https://www.nature.com/articles/s42256-018-0006-z.pdf>
- [13] So et al. The Evolved Transformer. *arXiv preprint arXiv:1901.11117*, 2019. <https://arxiv.org/abs/1901.11117>
- [14] Silver et al. Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm. *arXiv preprint arXiv:1712.01815*, 2017. <https://arxiv.org/abs/1712.01815>
- [15] Silver et al. AlphaGo Zero: Learning from scratch. In *DeepMind’s Blog*, 2017. <https://deepmind.com/blog/alphago-zero-learning-scratch/>
- [16] Andrychowicz et al. Learning to learn by gradient descent by gradient descent. *arXiv preprint arXiv:1606.04474*, 2016. <https://arxiv.org/abs/1606.04474>
- [17] Nichol et al. Reptile: A Scalable Meta-Learning Algorithm. 2018. <https://blog.openai.com/reptile/>
- [18] Frans et al. Meta Learning Shared Hierarchies. *arXiv preprint arXiv:1710.09767*, 2017. <https://arxiv.org/abs/1710.09767>
- [19] Zoph and Le, 2017 Neural Architecture Search with Reinforcement Learning. *arXiv preprint arXiv:1611.01578*, 2017. <https://arxiv.org/abs/1611.01578>
- [20] Salimans et al. Evolution Strategies as a Scalable Alternative to Reinforcement Learning. 2017. <https://blog.openai.com/evolution-strategies/>
- [21] Lehman et al. The Surprising Creativity of Digital Evolution: A Collection of Anecdotes from the Evolutionary Computation and Artificial Life Research Communities. *arXiv preprint arXiv:1803.03453*, 2018. <https://arxiv.org/abs/1803.03453>
- [22] Wang et al. Paired Open-Ended Trailblazer (POET): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions. *arXiv preprint arXiv:1901.01753*, 2019. <https://arxiv.org/abs/1901.01753>
- [23] Foerster et al. Learning to Model Other Minds. 2018. <https://blog.openai.com/learning-to-model-other-minds/>
- [24] Rabinowitz et al. Machine Theory of Mind. *arXiv preprint arXiv:1802.07740*, 2018. <https://arxiv.org/abs/1802.07740>